

Lecture „Neural Networks and Memristive Hardware Accelerators“
Prof. Dr. phil. nat. habil. Ronald Tetzlaff
Winter semester 2024/25

Lecture 5: PyTorch Tutorial

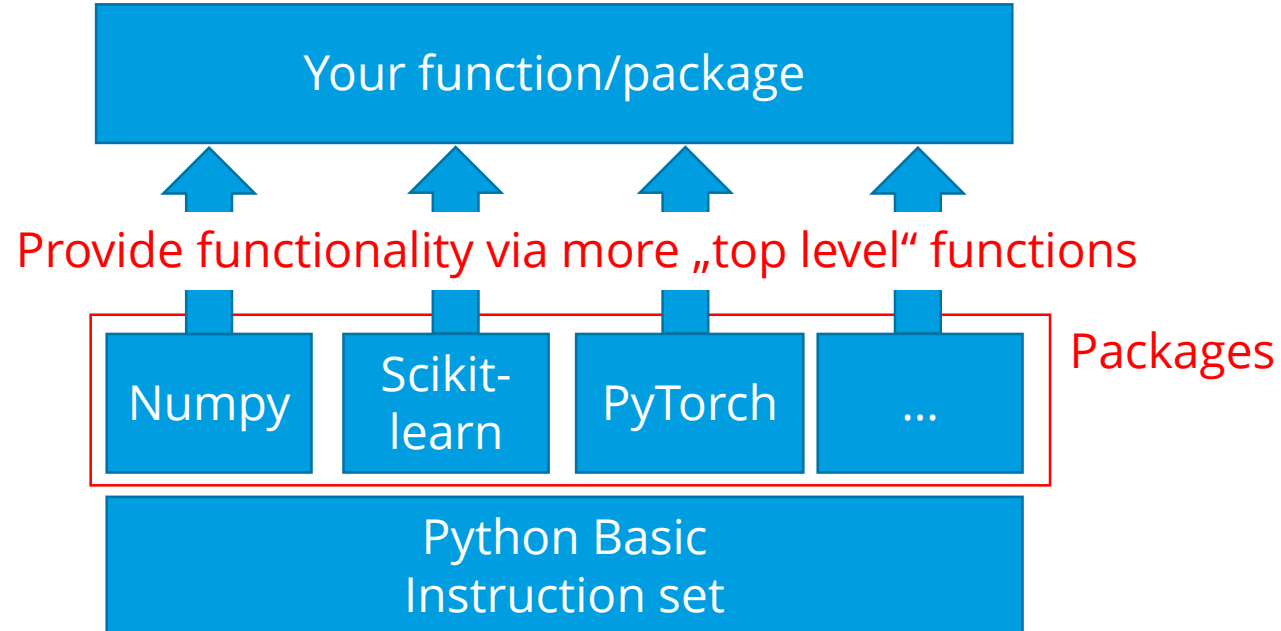
Dipl.-Ing. Steffen Seitz

PyTorch Tutorial : Study goals

- You understand what PyTorch can be used for.
- You know what a computation graph is.
- You are able to implement your own Neural Network architecture based on PyTorch.
- You know about training and evaluation mode in PyTorch.

PyTorch as a Module

Packages/Modules are nothing else than **classes** with **functions** written by other users free of charge! PyTorch is maintained by the „Meta“ AI Research Team.



What is PyTorch?

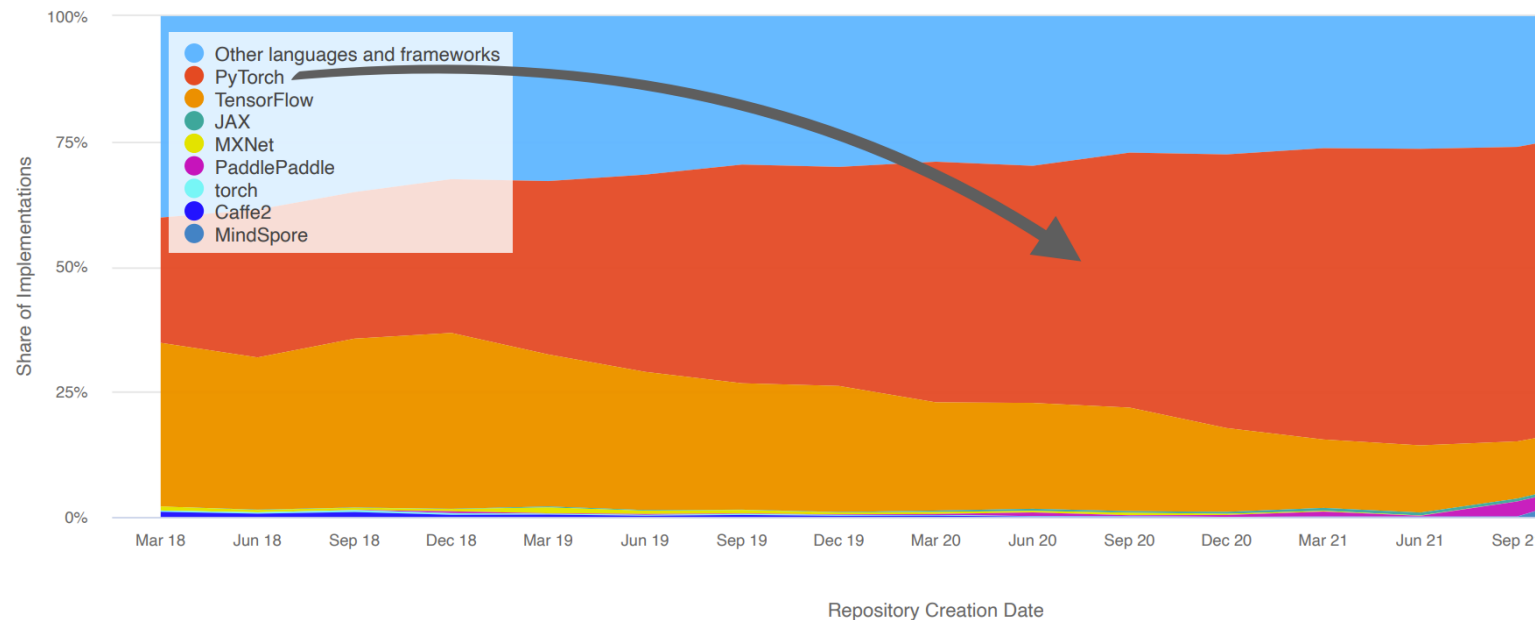
PyTorch is a scripting language that is very „pythonic“ in terms of readability.

Tensor
Library

Deep Learning
Library

Automatic
differentiation engine

Paper Implementations grouped by framework



PyTorch as a Tensor Computing Library

Scalar
(rank-0 tensor)

```
import torch

a = torch.tensor(1.)
a.shape

torch.Size([])
```

Vector
(rank-1 tensor)

```
a = torch.tensor([1., 2., 3.])
a.shape

torch.Size([3])
```

Matrix
(rank-2 tensor)

```
a = torch.tensor([[1., 2., 3.],
                  [2., 3., 4.]])
a.shape

torch.Size([2, 3])
```

`torch.tensor` \approx `numpy.array`

Color Images are a Stack of Matrices

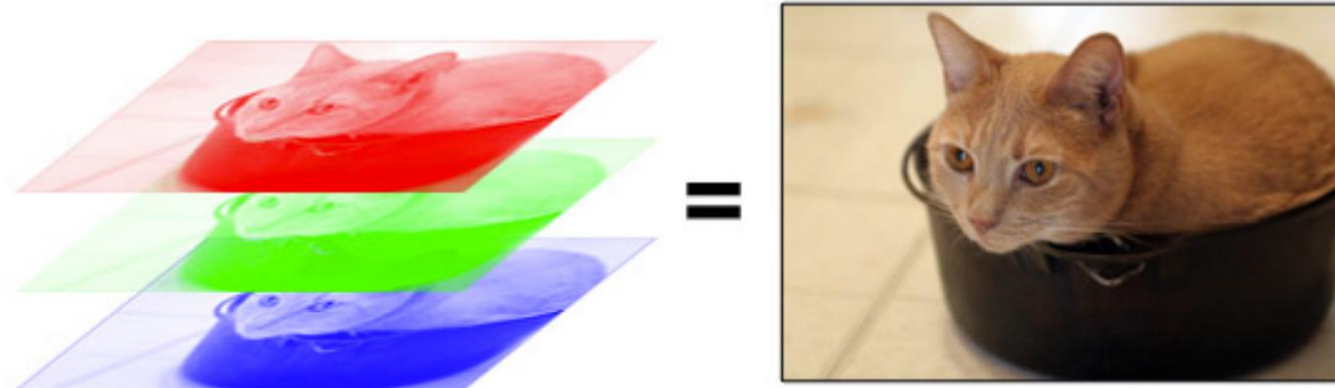


Image Source: <https://code.tutsplus.com/tutorials/create-a-retro-crt-distortion-effect-using-rgb-shifting--active-3359>

3D tensor
(rank-3 tensor)

```
a = torch.tensor([[[1., 2., 3.],  
                  [2., 3., 4.]],  
                  [[5., 6., 7.],  
                  [8., 9., 10.]])
```

```
a.shape
```

```
torch.Size([2, 2, 3])
```



Stack of Colour Images = 4D tensor
(rank-4 tensor)

```
b = torch.stack((a, a))  
b.shape
```

```
torch.Size([2, 2, 2, 3])
```

```
b
```

```
tensor([[[[ 1.,  2.,  3.],  
           [ 2.,  3.,  4.]],  
        [[ 5.,  6.,  7.],  
           [ 8.,  9., 10.]],
```

```
        [[ 1.,  2.,  3.],  
           [ 2.,  3.,  4.]],
```

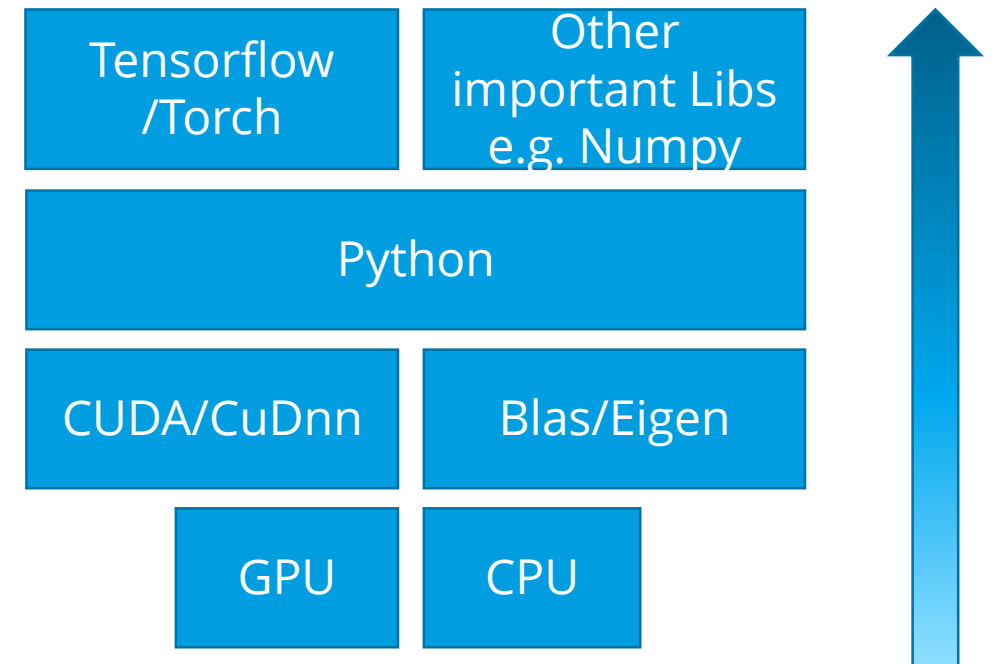
```
        [[ 5.,  6.,  7.],  
           [ 8.,  9., 10.]])])
```

PyTorch and CUDA

PyTorch has CUDA Support!

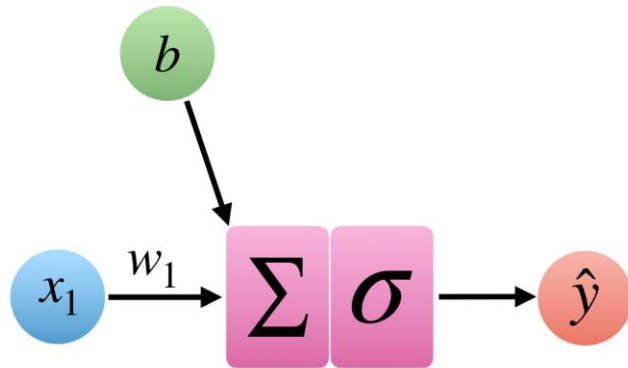
CUDA is a library by Nvidia used to speed up deep learning methods on GPUs.

It enables the transfer of data from the memory to the GPU accessible memory

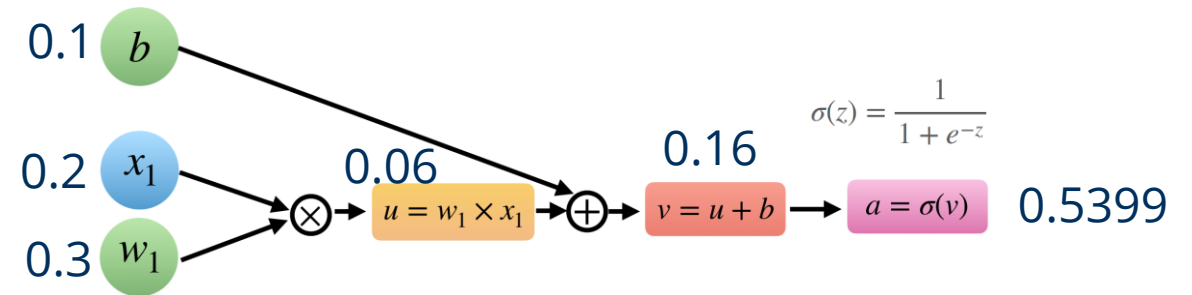


Computation Graph (Forward)

Perceptron (Forward Path)



Computational Graph (Forward Path)



Forward pass
in PyTorch

```
b = torch.tensor(0.1)
x1 = torch.tensor(0.2)
w1 = torch.tensor(0.3)
```

```
u = w1*x1
v = u + b
a = torch.sigmoid(v)
a
```

```
tensor(0.5399)
```


Computation Graph (Backward)

Forward pass in PyTorch

```
b = torch.tensor(0.1)
x1 = torch.tensor(0.2)
w1 = torch.tensor(0.3)
```

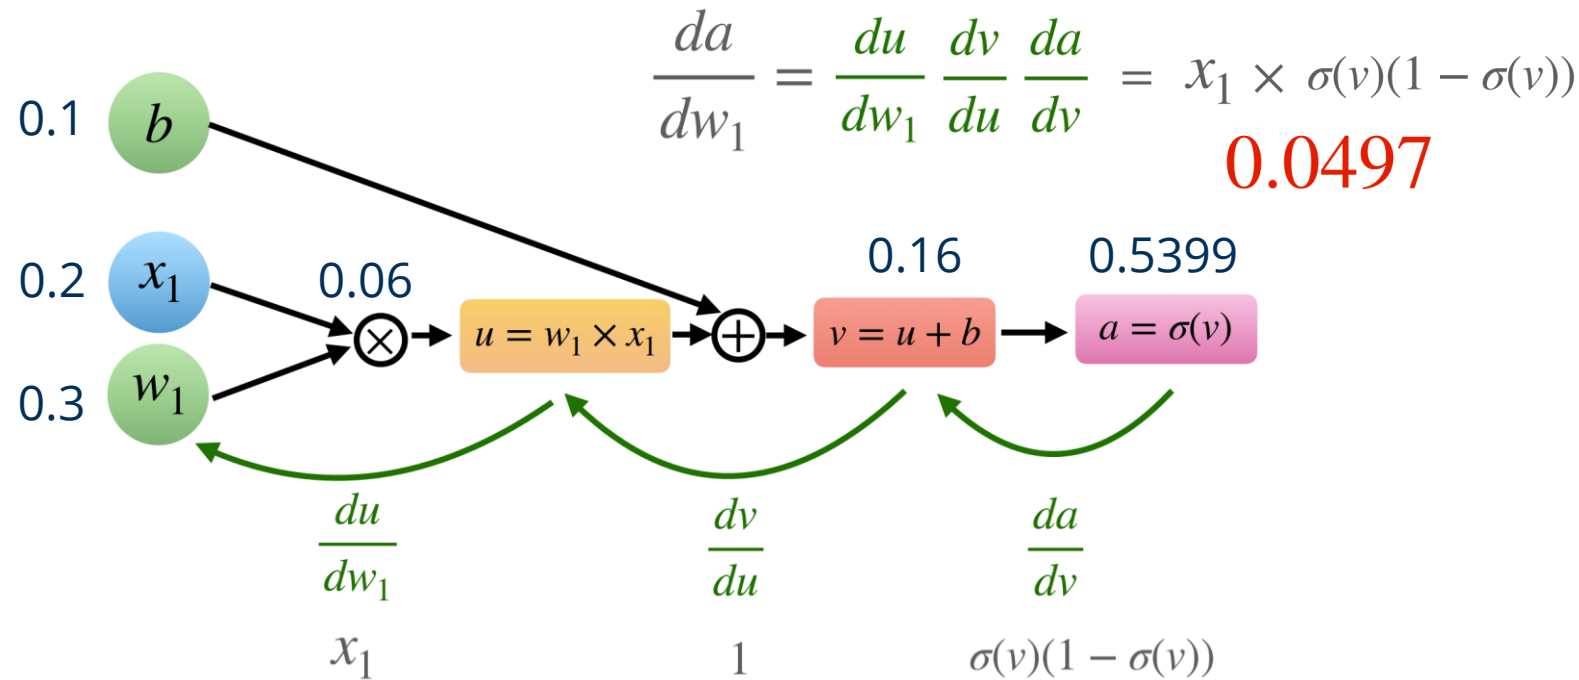
```
u = w1*x1
v = u + b
a = torch.sigmoid(v)
a
```

tensor(0.5399)

Backward pass in PyTorch

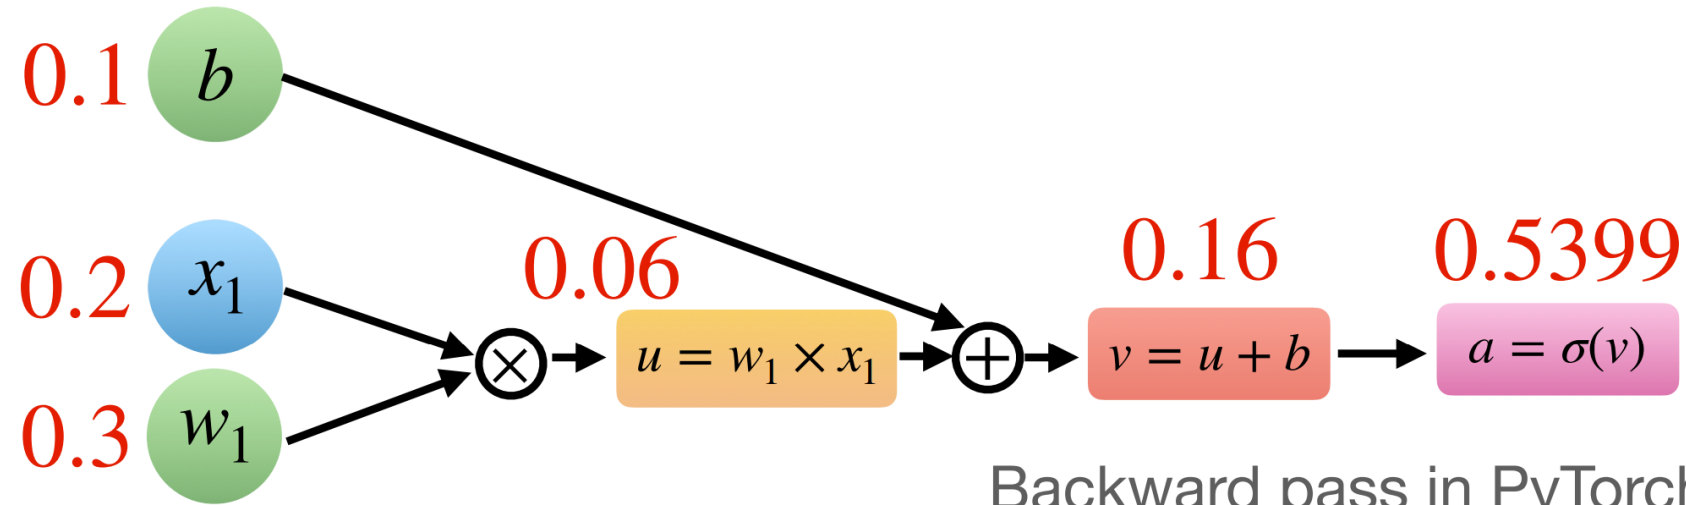
$$\frac{da}{dw_1} = a * (1-a) * x1$$

tensor(0.0497)



It would be nice if we would be able to derive this automatically! → **PyTorch Autograd!**

Autograd



```
b = torch.tensor(0.1)
x1 = torch.tensor(0.2)
w1 = torch.tensor(0.3, requires_grad=True)

u = w1*x1
v = u + b
a = torch.sigmoid(v)
a

tensor(0.5399, grad_fn=<SigmoidBackward0>)
```

Backward pass in PyTorch
automatically!

$$\frac{da}{dw_1} =$$

```
from torch.autograd import grad
grad(a, w1)

(tensor(0.0497),)
```

Exercise „Neural Network Implementation using Numpy“

Task: Build a Neural Network with a single hidden layer by only using Numpy!

Steps:

1. **Go** to the pre-prepared Colab: https://colab.research.google.com/drive/1d_00iXxgvC1kztQilPVkPsQAF4k1Bzf
2. **Clone the code** to your personal space, so you can change it.
3. **Follow the instructions** in the notebook and fill out the gaps with your own solution!
4. **Change the dataset** creation function from „moons“ to „circles“ or „blob“

Deep Learning with PyTorch - Example

Training images



<https://colab.research.google.com/drive/10696qeq3kyeibe2LD3UwkMXY3iZr-oN2?usp=sharing>

Deep Learning Workflow in Python

Import all dependencies (e.g. Pytorch)

```
import pandas as pd
import torch, torchvision
from torch import nn, optim
import torch.nn.functional as F
from matplotlib import pyplot as plt
```

Set Hyperparameters / Define GPU

```
BATCH_SIZE = 64;
EPOCHS = 3;
LEARNING_RATE = 0.001;
RANDOM_SEED = 1
# GPU Device
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

Import Dataset (e.g. Pandas)

```
pd.read_csv('data.csv')
```

If your CSV is in your working directory

Dataset Preprocessing (Next Slide)

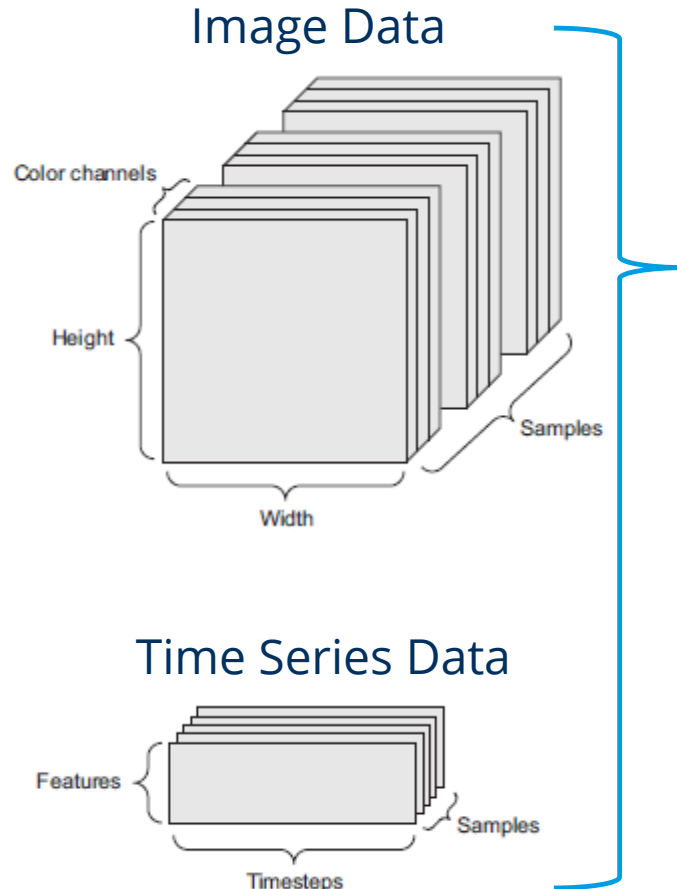
Define the Model

Define the Training loop

Train and Evaluate!

PyTorch

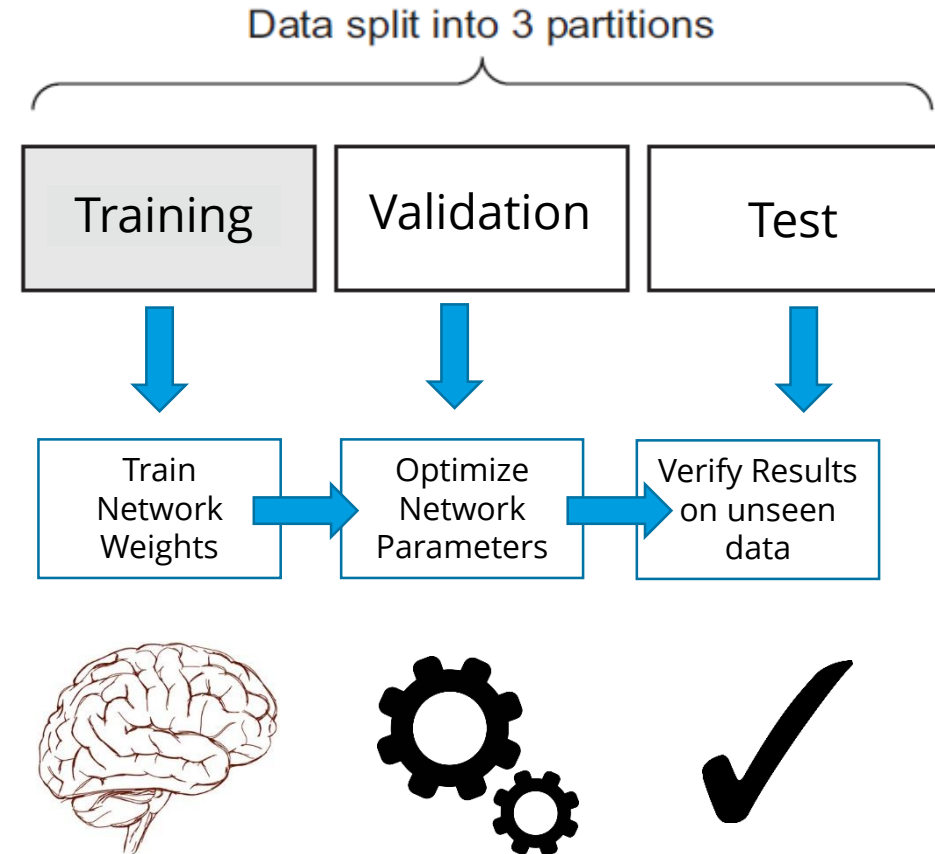
Data Preprocessing



Normalization to counteract skewed data/classes:

Normalization

$$\mathbf{x}_{j,std} = \frac{\mathbf{x}_j - \mu_j}{\sigma_j}$$



Define the Model

```
import pandas as pd
import torch, torchvision
from torch import nn, optim
import torch.nn.functional as F
from matplotlib import pyplot as plt
```

```
class FeedForwardNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(784, 1024); #fc1 = fully connected 1
        self.fc2 = nn.Linear(1024, 512);
        self.fc3 = nn.Linear(512, 256);
        self.out = nn.Linear(256, 10);

    def forward(self, X):
        # Flatten X for it to be able to be passed through a linear layer
        # Shape of X will be [64, 1, 28, 28] before the view() call and [64, 784] after
        X = X.view(X.shape[0], -1);

        # Pass X through first, second, third and output linear layer
        X = F.relu(self.fc1(X));
        X = F.relu(self.fc2(X));
        X = F.relu(self.fc3(X));
        X = self.out(X);

        return X;
```

Layer with a
state go here

Define how things are
executed

Initialize the Model and Define the Training Loop

Check if a GPU is available

```
device  
device(type='cuda', index=0)
```

We can ask for the device variable to check our device (GPU or CPU)

```
torch.manual_seed(RANDOM_SEED)
```

Additionally we save our random seed for reproducibility

```
model = FeedForwardNetwork();  
'If you want to use GPU'  
# model = model.to(device)  
criterion = nn.CrossEntropyLoss();  
optimizer = optim.Adam(model.parameters(), lr=LEARNING_RATE);
```

Now we create the instance from our model class and define the loss function and optimizer

```
model  
  
FeedForwardNetwork(  
  (fc1): Linear(in_features=784, out_features=1024, bias=True)  
  (fc2): Linear(in_features=1024, out_features=512, bias=True)  
  (fc3): Linear(in_features=512, out_features=256, bias=True)  
  (out): Linear(in_features=256, out_features=10, bias=True)  
)
```

Our model looks like expected

Initialize the Model and Define the Training Loop

```
def validate(val_batch):
    model.eval();
    with torch.no_grad():
        X, y = val_batch;
        out = model(X);

        predictions = torch.argmax(out, dim=1);
        samples_correct = predictions[predictions == y].shape[0];

        val_loss = criterion(out, y); ' criterion = nn.CrossEntropyLoss(); '
        val_accuracy = samples_correct/X.shape[0] * 100;

    model.train();
    return val_loss, val_accuracy;
```

```
class FeedForwardNetwork(nn.Module):
    def __init__(self):
        super().__init__();
        self.fc1 = nn.Linear(784, 1024);      #fc1 = fully connected 1
        self.fc2 = nn.Linear(1024, 512);
        self.fc3 = nn.Linear(512, 256);
        self.out = nn.Linear(256, 10);

    def forward(self, X):
        # Flatten X for it to be able to be passed through a linear layer
        # Shape of X will be [64, 1, 28, 28] before the view() call and [64, 784] after
        X = X.view(X.shape[0], -1);

        # Pass X through first, second, third and output linear layer
        X = F.relu(self.fc1(X));
        X = F.relu(self.fc2(X));
        X = F.relu(self.fc3(X));
        X = self.out(X);

        return X;
```

Last layer seems to miss a softmax which is needed for classification

Softmax is built into the CrossEntropyLoss in PyTorch! Therefore is not needed here!

The Training Loop

If you omit this, gradients will be accumulated, which we usually don't want

Softmax is applied

The gradients are calculated

The optimizer updates the weights

Here we track the performance
(Optional but recommended)

```
train_losses = [];  
train_accuracies = [];  
val_losses = [];  
val_accuracies = [];  
  
val_iterator = iter(val_loader);  
  
for epoch in range(EPOCHS):  
    model.train()  
    for i, (X, y) in enumerate(train_loader):  
        # X is a tensor with shape [64, 1, 28, 28] (the batch of images)  
        # 64: batchsize, 1: image channels (1 because grayscale), 28: image height, 28: image width  
        # y is a tensor with shape [64] (the batch of labels)  
        optimizer.zero_grad();  
  
        out = model(X);  
        # out (tensor): [64, 10]  
  
        predictions = torch.argmax(out, dim=1);  
        batch_samples_correct = predictions[predictions == y].shape[0];  
        batch_accuracy = batch_samples_correct / y.shape[0] * 100;  
        batch_loss = criterion(out, y);  
  
        batch_loss.backward();  
        optimizer.step();  
  
        try:  
            val_loss, val_accuracy = validate(val_iterator.next());  
        except StopIteration:  
            val_iterator = iter(val_loader);  
            val_loss, val_accuracy = validate(val_iterator.next());  
  
        train_losses.append(batch_loss.item());  
        train_accuracies.append(batch_accuracy);  
        val_losses.append(val_loss.item());  
        val_accuracies.append(val_accuracy);
```

Feed Forward is performed

Evaluation/Training Mode

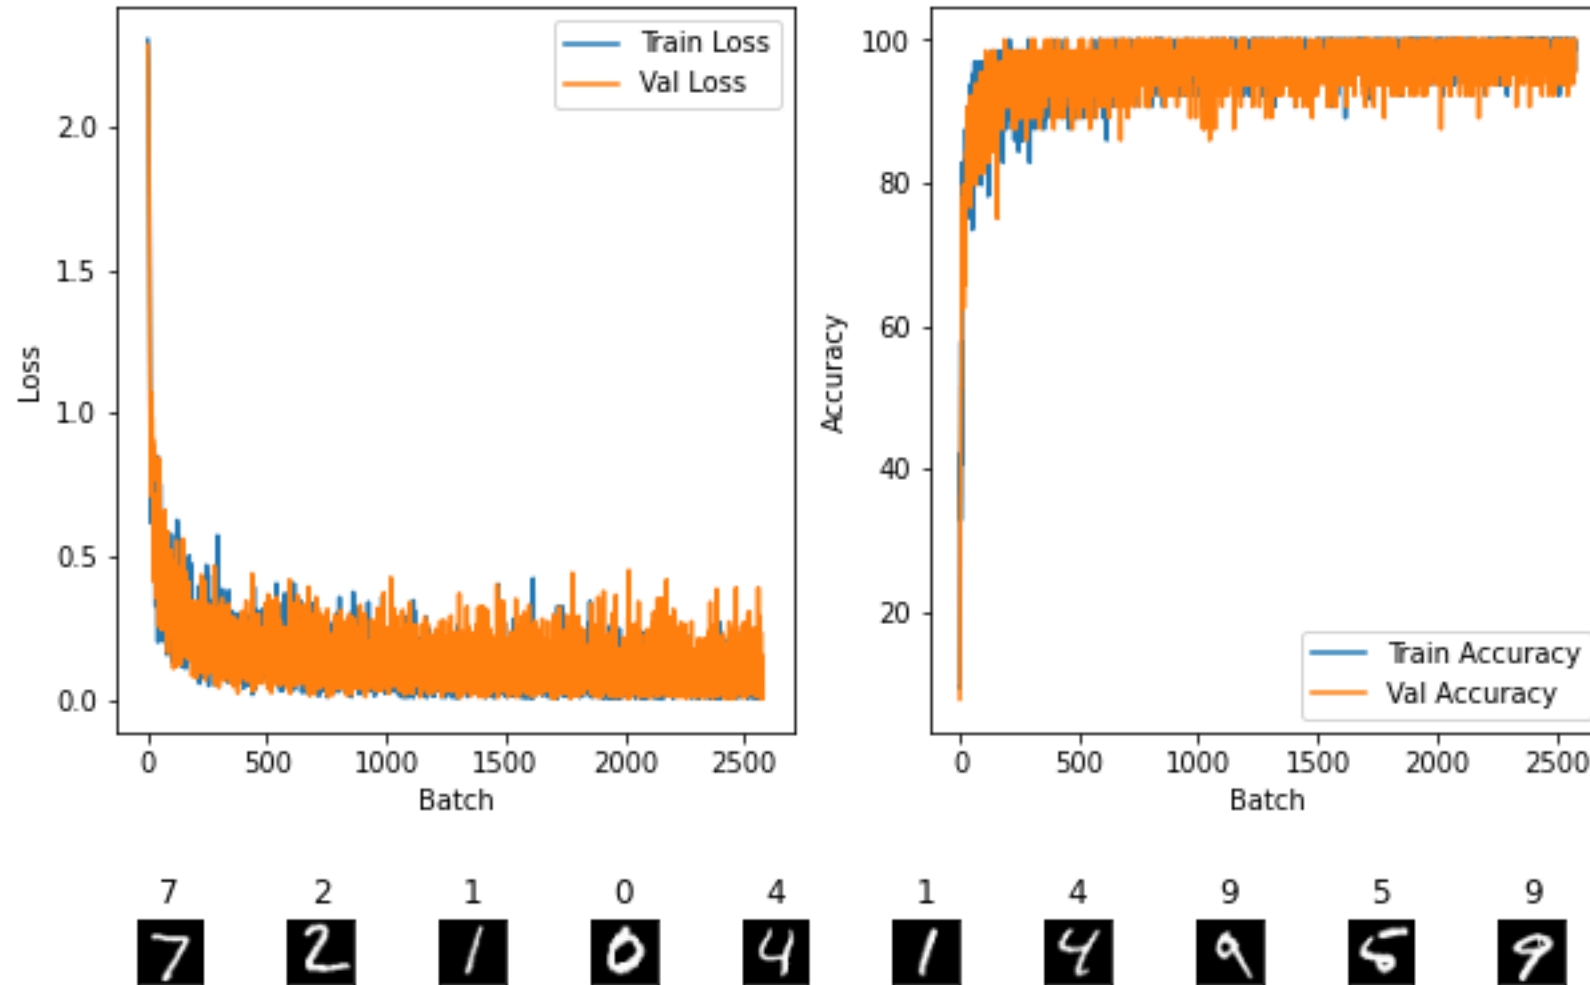
train() & .eval() modes matter for things like BatchNorm, Dropout etc.,

no.grad() prevents unnecessary graph construction

```
def validate(val_batch):  
    model.eval();  
    with torch.no_grad():  
        X, y = val_batch;  
        out = model(X);  
  
        predictions = torch.argmax(out, dim=1);  
        samples_correct = predictions[predictions == y].shape[0];  
  
        val_loss = criterion(out, y); ' criterion = nn.CrossEntropyLoss();  
        val_accuracy = samples_correct/X.shape[0] * 100;  
  
    model.train();  
    return val_loss, val_accuracy;
```

```
train_losses = [];  
train_accuracies = [];  
val_losses = [];  
val_accuracies = [];  
  
val_iterator = iter(val_loader);  
  
for epoch in range(EPOCHS):  
    model.train()  
    for i, (X, y) in enumerate(train_loader):  
        # X is a tensor with shape [64, 1, 28, 28] (the batch of images)  
        # 64: batchsize, 1: image channels (1 because grayscale), 28: image height, 28: image width  
        # y is a tensor with shape [64] (the batch of labels)  
        optimizer.zero_grad();  
  
        out = model(X);  
        # out (tensor): [64, 10]  
  
        predictions = torch.argmax(out, dim=1);  
        samples_correct = predictions[predictions == y].shape[0];  
        accuracy = batch_samples_correct / y.shape[0] * 100;  
        loss = criterion(out, y);  
  
        loss.backward();  
        zer.step();  
  
        l_loss, val_accuracy = validate(val_iterator.next());  
        StopIteration:  
        l_iterator = iter(val_loader);  
        l_loss, val_accuracy = validate(val_iterator.next());  
  
        train_losses.append(batch_loss.item());  
        train_accuracies.append(batch_accuracy);  
        val_losses.append(val_loss.item());  
        val_accuracies.append(val_accuracy);
```

Plot the Loss Functions and Accuracy over Time



Deep Learning with PyTorch – Live Demo



<https://colab.research.google.com/drive/1IRwU6IEWwMW2oxNAWzE14elkSr8q3792>

Literature

Code examples

- MNIST: <https://github.com/rasbt/stat453-deep-learning-ss21/blob/main/L13/code/1-lenet5-mnist.ipynb>
- CIFAR10: <https://github.com/rasbt/stat453-deep-learning-ss21/blob/main/L13/code/3-cnn-cifar10.ipynb>

Books

- Eli Stevens et al. „Deep Learning with PyTorch“, 2020, ISBN 9781617295263
https://isip.piconepress.com/courses/temple/ece_4822/resources/books/Deep-Learning-with-PyTorch.pdf
- Sebastian Raschka “Machine Learning with PyTorch and Scikit-Learn” 2022, ISBN: 978-1801819312
<https://learning.oreilly.com/library/view/machine-learning-with/9781801819312/> - Login via SLUB account

Next lecture: 6. Sequential Neural Networks (Steffen Seitz)